

```

In [1]: import torch
        from torch import nn
        from d2l import torch as d2l

In [2]: batch_size = 256 # 指定每个小批量包含的样本数量, 这里设置为256
        train_iter, test_iter = d2l.load_data_fashion_mnist(batch_size)
        # load_data_fashion_mnist() 函数, 数据集被自动下载并存储在默认的本地数据集目录下
        # train_iter 和 test_iter 分别表示训练集和测试集的迭代器, 可以用于获取小批量的数据样本和标签。

In [3]: # 单层隐藏层
        num_inputs, num_outputs, num_hiddens = 784, 10, 256 # num_inputs 和 num_outputs 分别是输入特征数和输出个数, num_hiddens 是隐
        # [Fashion-MNIST 数据集中每个样本都是 28x28 像素的灰度图像, 展开后共有 784 个特征, 而输出层需要输出 10 个类别的概率]

        W1 = nn.Parameter(torch.randn(num_inputs, num_hiddens, requires_grad=True) * 0.01) # W1 是隐藏层的权重矩阵
        b1 = nn.Parameter(torch.zeros(num_hiddens, requires_grad=True)) # b1 是隐藏层的偏置向量
        W2 = nn.Parameter(torch.randn(num_hiddens, num_outputs, requires_grad=True) * 0.01) # W2 是输出层的权重矩阵
        b2 = nn.Parameter(torch.zeros(num_outputs, requires_grad=True)) # b2 是输出层的偏置向量
        params = [W1, b1, W2, b2]

In [4]: # ReLU 激活函数
        def relu(X):
            a = torch.zeros_like(X) # 创建一个形状和 X 相同的全零张量 a
            return torch.max(X, a) # 计算输入 X 和全零张量 a 逐元素的最大值

```

```


In [5]: # 具有单层隐藏层的多层感知机模型
        # 输入 X 通过隐藏层的权重矩阵 W1 和偏置向量 b1 进行线性变换, 并经过 ReLU 激活函数的作用, 得到隐藏层的输出 H
        def net(X):
            X = X.reshape((-1, num_inputs))
            # 为了方便矩阵乘法的计算, 我们需要将其重塑为形状为 (batch_size, num_inputs) 的张量
            # 将其中一个维度设为 -1, 表示让 PyTorch 自动计算该维度的长度
            H = relu(X@W1 + b1) # 这里 "@" 代表矩阵乘法
            return (H@W2 + b2) # 对隐藏层的输出 H 进行线性变换, 返回输出层的输出张量

In [6]: loss = nn.CrossEntropyLoss(reduction='none') # 交叉熵损失函数, 模型的预测结果与真实标签之间的差距作为损失值

In [7]: num_epochs, lr = 10, 0.1
        updater = torch.optim.SGD(params, lr=lr)
        d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, updater)
        d2l.plt.title('1 hidden')

Out[7]: Text(0.5, 1.0, '1 hidden')

```




```

In [4]: # 更改学习率 0.1 -> 0.5
        num_inputs, num_outputs, num_hiddens = 784, 10, 256
        W1 = nn.Parameter(torch.randn(num_inputs, num_hiddens, requires_grad=True) * 0.01)
        b1 = nn.Parameter(torch.zeros(num_hiddens, requires_grad=True))
        W2 = nn.Parameter(torch.randn(num_hiddens, num_outputs, requires_grad=True) * 0.01)
        b2 = nn.Parameter(torch.zeros(num_outputs, requires_grad=True))
        params = [W1, b1, W2, b2]
        def relu(X):
            a = torch.zeros_like(X)
            return torch.max(X, a)

        def net(X):
            X = X.reshape((-1, num_inputs))
            H = relu(X@W1 + b1) # 这里 "@" 代表矩阵乘法
            return (H@W2 + b2)
        loss = nn.CrossEntropyLoss(reduction='none')
        num_epochs, lr = 10, 0.5
        updater = torch.optim.SGD(params, lr=lr)
        d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, updater)
        d2l.plt.title('lr changed')

Out[4]: Text(0.5, 1.0, 'lr changed')

```




```

In [5]: # 更改学习率0.1->0.05
num_inputs, num_outputs, num_hiddens = 784, 10, 256
W1 = nn.Parameter(torch.randn(num_inputs, num_hiddens, requires_grad=True) * 0.01)
b1 = nn.Parameter(torch.zeros(num_hiddens, requires_grad=True))
W2 = nn.Parameter(torch.randn(num_hiddens, num_outputs, requires_grad=True) * 0.01)
b2 = nn.Parameter(torch.zeros(num_outputs, requires_grad=True))
params = [W1, b1, W2, b2]
def relu(X):
    a = torch.zeros_like(X)
    return torch.max(X, a)

def net(X):
    X = X.reshape((-1, num_inputs))
    H = relu(X@W1 + b1) # 这里"@"代表矩阵乘法
    return (H@W2 + b2)
loss = nn.CrossEntropyLoss(reduction='none')
num_epochs, lr = 10, 0.05
updater = torch.optim.SGD(params, lr=lr)
d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, updater)
d2l.plt.title('lr changed')

```

Out[5]: Text(0.5, 1.0, 'lr changed')




```

In [7]: # 更改单层隐藏层中的隐藏单元数256->256*2
num_inputs, num_outputs, num_hiddens = 784, 10, 256*2
W1 = nn.Parameter(torch.randn(num_inputs, num_hiddens, requires_grad=True) * 0.01)
b1 = nn.Parameter(torch.zeros(num_hiddens, requires_grad=True))
W2 = nn.Parameter(torch.randn(num_hiddens, num_outputs, requires_grad=True) * 0.01)
b2 = nn.Parameter(torch.zeros(num_outputs, requires_grad=True))
params = [W1, b1, W2, b2]
def relu(X):
    a = torch.zeros_like(X)
    return torch.max(X, a)

def net(X):
    X = X.reshape((-1, num_inputs))
    H = relu(X@W1 + b1) # 这里"@"代表矩阵乘法
    return (H@W2 + b2)
loss = nn.CrossEntropyLoss(reduction='none')
num_epochs, lr = 10, 0.1
updater = torch.optim.SGD(params, lr=lr)
d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, updater)
d2l.plt.title('1 hidden changed')

```

Out[7]: Text(0.5, 1.0, '1 hidden changed')



```

In [8]: # 更改单层隐藏层中的隐藏单元数256->128
num_inputs, num_outputs, num_hiddens = 784, 10, 128
W1 = nn.Parameter(torch.randn(num_inputs, num_hiddens, requires_grad=True) * 0.01)
b1 = nn.Parameter(torch.zeros(num_hiddens, requires_grad=True))
W2 = nn.Parameter(torch.randn(num_hiddens, num_outputs, requires_grad=True) * 0.01)
b2 = nn.Parameter(torch.zeros(num_outputs, requires_grad=True))
params = [W1, b1, W2, b2]


def relu(X):
    a = torch.zeros_like(X)
    return torch.max(X, a)

def net(X):
    X = X.reshape((-1, num_inputs))
    H = relu(X@W1 + b1) # 这里"@"代表矩阵乘法
    return (H@W2 + b2)

loss = nn.CrossEntropyLoss(reduction='none')
num_epochs, lr = 10, 0.1
updater = torch.optim.SGD(params, lr=lr)
d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, updater)
d2l.plt.title('1 hidden changed')

```

Out[8]: Text(0.5, 1.0, '1 hidden changed')



```

In [9]: # 更改隐藏层数1->2
num_inputs, num_outputs, num_hiddens1, num_hiddens2 = 784, 10, 256, 256
W1 = nn.Parameter(torch.randn(num_inputs, num_hiddens1, requires_grad=True) * 0.01)
b1 = nn.Parameter(torch.zeros(num_hiddens1, requires_grad=True))
W2 = nn.Parameter(torch.randn(num_hiddens1, num_hiddens2, requires_grad=True) * 0.01)
b2 = nn.Parameter(torch.zeros(num_hiddens2, requires_grad=True))
W3 = nn.Parameter(torch.randn(num_hiddens2, num_outputs, requires_grad=True) * 0.01)
b3 = nn.Parameter(torch.zeros(num_outputs, requires_grad=True))
params = [W1, b1, W2, b2, W3, b3]


def relu(X):
    a = torch.zeros_like(X)
    return torch.max(X, a)

def net(X):
    X = X.reshape((-1, num_inputs))
    H = relu(X@W1 + b1)
    return (H@W2 + b2)@W3 + b3

loss = nn.CrossEntropyLoss(reduction='none')
num_epochs, lr = 10, 0.1
updater = torch.optim.SGD(params, lr=lr)
d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, updater)
d2l.plt.title('2 hidden')

```

Out[9]: Text(0.5, 1.0, '2 hidden')



```


In [10]: # 更改隐藏层数1->3
num_inputs, num_outputs, num_hidden1, num_hidden2, num_hidden3 = 784, 10, 256, 256, 256
W1 = nn.Parameter(torch.randn(num_inputs, num_hidden1, requires_grad=True) * 0.01)
b1 = nn.Parameter(torch.zeros(num_hidden1, requires_grad=True))
W2 = nn.Parameter(torch.randn(num_hidden1, num_hidden2, requires_grad=True) * 0.01)
b2 = nn.Parameter(torch.zeros(num_hidden2, requires_grad=True))
W3 = nn.Parameter(torch.randn(num_hidden2, num_hidden3, requires_grad=True) * 0.01)
b3 = nn.Parameter(torch.zeros(num_hidden3, requires_grad=True))
W4 = nn.Parameter(torch.randn(num_hidden3, num_outputs, requires_grad=True) * 0.01)
b4 = nn.Parameter(torch.zeros(num_outputs, requires_grad=True))
params = [W1, b1, W2, b2, W3, b3, W4, b4]

def relu(X):
    a = torch.zeros_like(X)
    return torch.max(X, a)

def net(X):
    X = X.reshape((-1, num_inputs))
    H = relu(X@W1 + b1)
    return (H@W4 + b4)
loss = nn.CrossEntropyLoss(reduction='none')
num_epochs, lr = 10, 0.1
updater = torch.optim.SGD(params, lr=lr)
d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, updater)
d2l.plt.title('3 hidden')

```

Out[10]: Text(0.5, 1.0, '3 hidden')



```

In [11]: # 更改隐藏层数1->4
num_inputs, num_outputs, num_hidden1, num_hidden2, num_hidden3, num_hidden4 = 784, 10, 256, 256, 256, 256
W1 = nn.Parameter(torch.randn(num_inputs, num_hidden1, requires_grad=True) * 0.01)
b1 = nn.Parameter(torch.zeros(num_hidden1, requires_grad=True))
W2 = nn.Parameter(torch.randn(num_hidden1, num_hidden2, requires_grad=True) * 0.01)
b2 = nn.Parameter(torch.zeros(num_hidden2, requires_grad=True))
W3 = nn.Parameter(torch.randn(num_hidden2, num_hidden3, requires_grad=True) * 0.01)
b3 = nn.Parameter(torch.zeros(num_hidden3, requires_grad=True))
W4 = nn.Parameter(torch.randn(num_hidden3, num_hidden4, requires_grad=True) * 0.01)
b4 = nn.Parameter(torch.zeros(num_hidden4, requires_grad=True))
W5 = nn.Parameter(torch.randn(num_hidden4, num_outputs, requires_grad=True) * 0.01)
b5 = nn.Parameter(torch.zeros(num_outputs, requires_grad=True))
params = [W1, b1, W2, b2, W3, b3, W4, b4, W5, b5]

def relu(X):
    a = torch.zeros_like(X)
    return torch.max(X, a)

def net(X):
    X = X.reshape((-1, num_inputs))
    H = relu(X@W1 + b1)
    return (H@W5 + b5)
loss = nn.CrossEntropyLoss(reduction='none')
num_epochs, lr = 10, 0.1
updater = torch.optim.SGD(params, lr=lr)
d2l.train_ch3(net, train_iter, test_iter, loss, num_epochs, updater)
d2l.plt.title('4 hidden')

```

Out[11]: Text(0.5, 1.0, '4 hidden')

